

Manage Kubernetes Resources Using CLI

Goal: Learn to install the tools and manage basic Kubernetes resources using the `kubectl` CLI: create pods and deployments, scale, expose as a service.

Prerequisites

- Laptop/desktop with **at least 4 GB RAM** and internet access.
- One of these OS options: **Ubuntu/Debian Linux, macOS, or Windows 10/11**.
- Basic terminal or PowerShell familiarity.

Quick Checklist Before You Start

- Terminal (Linux/macOS) or PowerShell (Windows) ready.
- Admin/sudo access to install software.
- Docker and Minikube installed.

Install the Tools

We will use **Minikube** for the local Kubernetes cluster. It runs a single-node cluster inside Docker.

You need: **Docker (or Docker Desktop)**, **kubectl**, and **Minikube**.

A. Install Docker

Windows:

1. Download and install **Docker Desktop** from Docker's website.
2. Enable **WSL2 backend** during installation.
3. Start Docker Desktop and ensure it is running.

Verify Docker installation:

```
docker --version
```

B. Install kubectl (Kubernetes CLI)

Windows (using Chocolatey):

```
choco install kubernetes-cli
```

Verify installation:

```
kubectl version --client
```

C. Install and Start Minikube

Windows (using Chocolatey):

```
choco install minikube  
minikube start --driver=docker
```

Verify the cluster is running:

```
minikube status  
kubectl get nodes
```

If `kubectl get nodes` returns a node in **Ready** state — you are ready to go!

❖❖ Step 1: Create a Simple

Pod

Create a file: `nginx-pod.yaml`

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:latest
      ports:
        - containerPort: 80
```

Apply the manifest:

```
kubectl apply -f nginx-pod.yaml
```

Check pod status:

```
kubectl get pods
# expected: nginx-pod 1/1 Running
```

Describe the pod (to see detailed info):

```
kubectl describe pod nginx-pod
```

Forward the port locally to access it in the browser:

```
kubectl port-forward pod/nginx-pod 8080:80
```

Then open <http://localhost:8080> in your browser.

When done, delete the pod:

```
kubectl delete pod nginx-pod
```

◆◆ Step 2: Deployment and

Scaling Create a deployment:

```
kubectl create deployment my-nginx --image=nginx
```

Check deployments and pods:

```
kubectl get deployments  
kubectl get pods -l app=my-nginx
```

Scale to 3 replicas:

```
kubectl scale deployment my-nginx --replicas=3  
kubectl get pods
```

Update the image (rolling update):

```
kubectl set image deployment/my-nginx nginx=nginx:1.25  
kubectl rollout status deployment/my-nginx
```

If something goes wrong, rollback:

```
kubectl rollout undo deployment/my-nginx
```

◆◆ Step 3: Expose Deployment as a

Service Expose the deployment using a NodePort service:

```
kubectl expose deployment my-nginx --type=NodePort --port=80  
kubectl get svc
```

Get the URL of your app using Minikube:

```
minikube service my-nginx --url
```

Open the displayed URL in your browser.
To remove the service:

```
kubectl delete svc my-nginx
```

❖❖ Step 4: Cleanup

After the lab, clean up all created resources:

```
kubectl delete deployment my-nginx  
kubectl delete svc my-nginx  
kubectl delete all --all -n default
```

If you want to stop or delete the cluster:

```
minikube stop  
minikube delete
```

❖❖ Summary

In this lab, you learned how to:

1. Install Docker, kubectl, and Minikube.
2. Create and manage Pods.
3. Deploy and scale applications.
4. Expose Deployments as services.
5. Clean up resources safely.

Expose Deployments as services.

Clean up resources safely.